

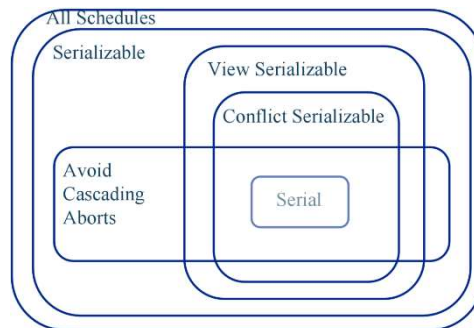
## CS186 Final Cheat Sheet

### Transactions & Currency Control:

- **DBMS:** database management system manages the queries
- **Concurrency control** provides correct and fast data access in the presence of concurrent work by many users
- **Recovery** ensures database is fault tolerant
- **Concurrent execution arguments:**
  - **Throughput** argument: increases TPS (transactions/second)
  - **Latency** argument: latency not dependent on another unrelated transaction
- Inconsistent/dirty reads and inconsistent updates could happen
- A **transaction** is a collection of operations that form a single logical unit (one atomic unit of work) (begin, SQL statements, end)
- **ACID:** *Atomicity*: all or none, *consistency*, starts and ends consistent, *isolation*: execute each Xact from others, *durability*: commit persists
- **Serial schedules:** each transaction runs from start to finish without any intervening actions from other transactions
- 2 schedules are **equivalent** if they involve same xact, each individual xact actions are ordered the same and leave the db in same final state
- **Serializable** if s is equivalent to any serial schedule
- **Conflict equivalent:** same actions, conflicts ordered same way
- **Conflict serializable:** can be equal to serial schedules. Steps: draw dependency graph. If cycle, not conflict serializable.
- **View serializable** = conflict serial + blind writes. Hard to enforce

### Locking

- **2PL** two phase locking is pessimistic, get all locks, then release all
  - Guarantees conflict serializability, doesn't prevent cascading aborts
  - **Strict 2PL** releases all locks at the same time
- **Lock manager** maintains a hashtable of locked objects, mode and wait queue
- Three ways of dealing with **deadlock**: prevention, avoidance, detection & resolution.
  - Prevention: use resource ordering (screen < network < print)



- Avoidance: Priorities based on age (now – start\_time)
- **Wait-Die:** if Ti has higher priority, Ti waits for Tj; else Ti aborts
- **Wound-Wait:** if Ti has higher priority, Tj aborts; else Ti waits
- Detection: generates “waits for” graphs, and checks if cycles

### Recovery

- Assume strict 2PL for recovery & in place updates
- **FORCE:** every update is on DB before commit
  - Durability without REDO, but poor perf
- **NO STEAL:** don't allow buffer-pool frames with uncommitted updates to be replaced
  - Atomicity without UNDO logging, poor perf
- **STEAL, NO-FORCE** is preferred policy

LogRecord fields:

LSN  
prevLSN  
XID  
type  
pageID  
length  
offset  
before-image  
after-image

update records only

	No Steal	Steal
No Force		Fastest
Force	Slowest	

Performance Implications

	No Steal	Steal
No Force	No UNDO REDO	UNDO REDO
Force	No UNDO No REDO	UNDO No REDO

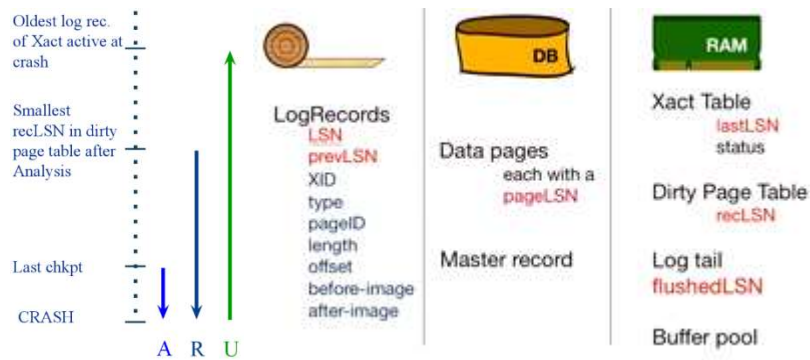
Logging/Recovery Implications

- **WAL (Write Ahead Logging)** is the protocol that we use
  - Force log record before data page (Atomicity)
  - Force all log records for an X before commit (durability)
- WAL Info: **LSN**: unique and increasing, **flushedLSN** (last logged LSN), **pageLSN**: points to an LSN from the DB. pageLSNi <= flushedLSN
- Possible log records: **update**, **commit**, **abort**, **checkpoint** (log maintenance), **CLRs** (compensation log records – never undone), **end** (end of commit or abort)
- Two in-memory tables for recovery:

Transaction Table		
XID	Status	lastLSN
1	R	33
2	C	42

Dirty Page Table	
PageID	recLSN
46	11
63	24

- Xact table: **lastLSN** (most recent LSN written)
- Dirty Page table: **recLSN** LSN that first caused the dirty page



- Assume disk write is atomic. Strict 2PL WAL.
- Checkpoints** have Xact and dirty page table
- 3 phases for crash recovery:
  - Analysis** figure out which Xacts committed since checkpoint, which failed.
  - REDO** all actions (repeat history)
  - UNDO** effects of failed Xacts
- Analysis:**
  - Reestablish knowledge of checkpoint state
  - End** recs: remove Xact from Xact table
  - Update** recs: if page P not in Dirty Page Table, Add P to DPT, set its recLSN=LSN
  - !End** recs: Add Xact to Xact table, set lastLST=LSN, change Xact status on commit
  - At end, for any Xacts in the Xact table in Committing state, generate a corresponding END log record, and remove Xact from Xact table
  - Xact table says which xacts were active at last log flush before crash
  - DPT says which dirty pages might not have made it to disk
- REDO:**
  - Reapply all updates (even aborted Xacts), redo CLR's.
  - Scan forward from log rec containing smallest recLSN in DPT
  - For each update log record or CLR with LSN, REDO action unless:
    - Affected page not in DPT
    - Affected page in DPT but has recLSN > LSN
    - pageLSN (in DB) >= LSN (IO required)
  - Reapply action, set pageLSN to LSN (no log, no force)
- UNDO:**

ToUndo={lastLSNs of all Xacts in the Xact Table} // a.k.a. "losers"

Repeat until ToUndo is empty:

- Choose (and remove) largest LSN among ToUndo.
- If this LSN is a CLR and undonextLSN==NULL
  - Write an End record for this Xact.
- If this LSN is a CLR, and undonextLSN != NULL
  - Add undonextLSN to ToUndo
- Else this LSN is an update. Undo the update, write a CLR, add prevLSN to ToUndo.

- Recovery Manager** guarantees Atomicity & Durability

### Replication Consistency and NoSQL

- Why replicate? **Doesn't help with data scaling** (this is what sharding is for), **does help with workload scaling** (load balancing)
- Replication increases availability (reduced latency)
- Single master replication:** both sides handle full update volume
- Single master log-shipping:** Log is shipped to secondary nodes
- Multi-master:** 2PC can cause high latency, no 2PC write conflicts
- Replication granularity can be at DB, table, partition or tuple level
- If coordinator fails: participants hold locks indefinitely (not good)
- Paxos** a consensus protocol like 2PC to get a group of nodes to agree on a proposal. Resilient to node failures (majority)
  - Replicate coordinator 2F + 1 replicas to tolerate F faults (virtually indestructible coordinator)

### Pre-Final Review

- Closure of attribute** is denoted with A+, says everything that can be determined by A.
- BCNF** (3NF & x->y, x is super key). Ex. A -> BCD, BC -> AD, D -> B. D->B isn't a super key (separate table to ADC and DB)
- Dependency preserving:** After decomposition, are the FD the same or not (do they fit with the new tables?)
- Lossless:** Perform alpha on tables.

