

## CS186 Midterm 2 Cheat Sheet

### Sorting and Hashing:

- Sorting # of passes =  $1 + \text{ceil}(\log_{B-1}(\text{ciel}(N/B)))$
- Sorting cost:  $2N * \text{\# of passes}$
- Hashing cost:  $2N * \text{\#passes} = 4N$  (assuming no overflow)
  - Divide (partitions to B-1)

### Relational Algebra:

- **Codd's theorem** established equivalence in expressivity between **relational calculus** and **relational algebra**
- **Unary Operators:**
  - **Projection ( $\pi$ ):** Retains only the desired columns (vertical)
  - **Selection ( $\sigma$ ):** Selects a subset of rows (horizontal) (WHERE)
  - **Renaming ( $\rho$ ):** Rename attributes and relations

$\rho(\text{Temp1}(1 \rightarrow \text{sid1}, 4 \rightarrow \text{sid2}), R1 \times S1)$

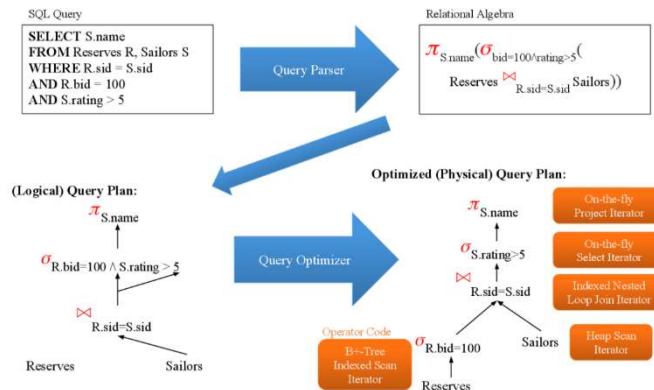
Output Relation Name      Renaming List position  $\rightarrow$  New Name      Input Relation

- **Binary Operators** (on pairs of relations):
  - **Union ( $\cup$ ):** Tuples in r1 or in r2 (removes duplicates)
  - **Set-Difference ( $-$ ):** Tuples in r1, but not in r2 (EXCEPT)
  - **Cross-Product ( $\times$ ):** Allows us to combine two relations.
- **Compound Operators** (common “macros” for the above):
  - **Intersection ( $\cap$ ):** Tuples in r1 and in r2
  - **Joins ( $\bowtie$ ,  $\Join$ ):** Combine relations that satisfy predicates
- **Division operator:** all entries  $x$  in  $A$  such that for all  $y$  in  $B$  there is an  $(x, y)$  in  $A$ .  $A / B = \pi_x(A) - \pi_x((\pi_x(A) \times B) - A)$
- **Group By/Aggregation operator ( $\gamma$ ):** (HAVING)  $\gamma_{age, \text{AVG}(rating), \text{COUNT(*)} > 2}(\text{Sailors})$  or GROUP BY

### Iterators and Joins

- **Cost Notation**
  - $|R|$ : the number of pages to store R

Index I on primary key matches selection:  $\text{Height}(I) + 1$   
 Clustered on one or more:  $(N\text{Pages}(I) + N\text{Pages}(R)) * \text{product of RF's of matching selects}$   
 Non-clustered index I matching one or more selections:  $(N\text{Pages}(I) + N\text{Tuples}^*) * \text{product of RF's of matching selects}$   
 Sequential scan:  $N\text{Pages}(R)$



- $p_R$ : number of records per page of R
- $|R|$ : the cardinality (number of records) of R
  - $|R| = p_R * [R]$
- **SNLJ:** foreach **record** r in R do for each **record** s in S do join
  - Cost:  $(p_R * [R]) * [S] + [R]$  (works for more than equality)
- **PNLJ:** Assume we use a small amount of RAM (to store a page of R and a page of S)
  - Foreach **page**  $b_R$  in R do: Foreach **page**  $b_S$  in S do:
  - Foreach **record** r in  $b_R$  do: Foreach **record** s in  $b_S$  do: join
  - Cost:  $[R] * [S] + [R]$
- **BNLJ:** Assume we use a B RAM (to store a page of R and a page of S)
  - Foreach **block**  $b_R$  of B-2 pages in R do: Foreach **page**  $b_S$  in S do:
  - Foreach **record** r in  $b_R$  do: Foreach **record** s in  $b_S$  do: join
  - Cost:  $\text{ciel}([R]/(B-2)) * [S] + [R]$
- **Index Nested Loops Join:** uses index lookups (on one of the inputs)
  - Foreach tuple r in R do find matching S tuples: join
  - Cost =  $[R] + ([R] * p_R) * \text{cost to find matching S tuples}$
- **Sort-Merge Join:** requires equality predicate (**equi-joins** and **natural joins**) 2 steps: **sort** and then **merge**
  - Cost: Sort R + Sort S +  $([R] + [S])$  (worst cast  $[R] * [S]$ )
- **Simple Hash Join:**
  - Same as INLJ on hash
  - Simple: R fits in memory
- **Grace Hash join:** requires equality predicate: **divide** and **conquer**
  - Partition tuples from R and S by join key  $(2([R] + [S]) \text{ I/Os})$
  - Matching:  $[R] + [S]$


```
while not done {
  while (r < s) { advance r }
  while (r > s) { advance s }
  // assert r == s
  mark s // save start of "block"
  while (r == s) {
    // Outer loop over r
    while (r == s) {
      // Inner loop over s
      yield <r, s>
      advance s
    }
    reset s to mark
    advance r
  }
}
```

### Parallel Query Processing:

- **Speed-up:** fixed workload, want linear performance to parallelism
- **Scale-up:** increase workload, want constant performance
- **Pipeline parallelism** and **partition parallelism**
- Architectures: shared **memory &**, shared **disk**, shared **nothing**
- Query parallelism:
  - **Inter-query** (parallelism across queries). 1 query/processor
    - Careful of concurrency control (easy)
  - **Intra-query** (between operators within a single query)
    - **Inter-operator** can allow for pipeline parallelism

**Armstrong's Axioms:** Reflexivity, augmentation, transitivity, union, decomposition.  
**Attribute closure:** all possibilities  
 1<sup>st</sup> normal form, all attribute atomic. **BCNF**  
 $X \rightarrow A$  in F+ if x is a superkey of R and A in X.

- $|f(x)| \mid |g(f(x))| \mid |h(g(f(x)))|$
- **Bushy** tree parallelism (left and right branch runs independently on 2 different branches)
- **Intra-operator** allows for between operators
- **Partition parallelism** stuff is partitioned to different machines
- **Data partitioning** is how we split data up across disks/machines
  - **Range:** like histograms. A-E in 1 bucket... (good for equijoins, range queries, group-by)
  - **Hash:** Each machine takes in the hash (equijoins/group-by)
  - **Round Robin:** As it implies (equally spaced load spreading)
- **Parallel hashing** is dead simple, add a step before first hash. Hash data to however many computers we plan to parallelize.
- **Parallel Grace Hash Join:**
  - Also dead simple. Perfect scale/speedup
- **Parallel sorting:** partition data by range (sample data) and data is sorted over 3 machines
- **Parallel sort merge join:** pass 0 ... n-1 are like parallel sorting. Range partition R & S + before
- **Join shuffling:**
  - **Asymmetric shuffle** is just when R is already partitioned
  - **Broadcast join** take a small S, and copy it to all of the R machines and join there



### Relational Query Optimization:



- **Query parser** checks correctness, authorization, generates parse tree
- **Query rewriter** converts queries to canonical form.
- **Plan space:** for a given query, what plans are considered
- **Cost estimation:** how is the cost of plan estimated
- **Search strategy:** how do we "search" in the "plan space?"
- Algebra equivalences:  $\sigma_{c1 \wedge \dots \wedge c_n}(R) \equiv \sigma_{c1}(\dots(\sigma_{c_n}(R))\dots)$  (cascade)
- $\sigma_{c1}(\sigma_{c2}(R)) \equiv \sigma_{c2}(\sigma_{c1}(R))$  (commute) ||  $\pi_{a1}(R) \equiv \pi_{a1}(\dots(\pi_{a1}, \dots, a_{n-1}(R))\dots)$  (cascade) ||  $R \times (S \times T) \equiv (R \times S) \times T$  (associative) ||  $R \times S \equiv S \times R$  (commutative)
- Apply **selections**, **projects** ASAP and avoid **Cartesian** products
- System R/Selinger optimizers is what we'll do. Cascades is other
  - Works well for 10-15 joins
  - Plan space: avoid overpriced subtree. Too large, must be pruned, only left-deep plans (avoid cartesian products)

- Cost estimation: inexact, but works ok
- Search algorithm: dynamic programming
- **Eager projection** you can project anything not needed later
- $\sigma_{R.a=S.b}(R \times S) \equiv R \bowtie_{R.a=S.b} S \mid \mid \sigma_{(R \bowtie S)} \equiv \sigma_{(R)} \bowtie S$
- Common physical properties of an output: sort order, hash grouping
- Index scan (sorted result), sort (sorted result), hash (grouped result)
- MergeJoin (input sorted, output sorted), NLJ (preserves sort order)
- System R: cost = #I/O + CPU-factor \* #tuples
- Selectivity = |output| / |input| = RF
- Size estimation (1/10 default):
  - Col=value (1/NKeys)
  - Col1=col2 (1/max(NKeys(1), NKeys(2)))
  - Col > value (high(l)-value)/(High(l)-Low(l) +1)
- Independence assumption:  $\sigma_p > 89 \wedge \sigma_{age < 26}$ ? 50% \* 46% = 23%
- $\sigma_p > 89 \vee \sigma_{age < 26}$ ? 50% + 46% - (50% x 46%) = 73%
- Database Design: requirements analysis, conceptual design, logical design, schema refinement, physical design, security design
- Data independence? **Hellerstein's Inequality:** If your env changes much faster than application, you need data independence.

Statistic	Meaning
NTuples	# of tuples in a table (cardinality)
NPages	# of disk pages in a table
Low/High	min/max value in a column
Nkeys	# of distinct values in a column
IHeight	the height of an index
INPages	# of disk pages in an index

### Data Models

- Entity set (rectangles), relationship set (diamond), attribute (oval)
- **Key constraints:** arrow, each dept has at most 1 manager
- **Participation constraints:** At least 1. Every employee works in a dept
- **Weak entity** can be identified uniquely by id of another entity



- **x->y** means x determines y
- **Superkey** a set of columns in a table that determines all.
- **Candidate key:** a minimal set of columns that determine all.
- **Primary key:** a single chosen candidate key
- **Update anomaly:** can't update just 1 column. **Insertion anomaly:** insert, but don't know wage, **deletion anomaly,** delete all employees with certain hour, we lose the info about wage